

An Analytic Manipulation Bound for the Carry-Over Lottery Allocation Family

Timothy Highley¹, Kaushik Rajan²

¹La Salle University

²Independent researcher

highley@lasalle.edu, kaushi@alumni.ncsu.edu

Abstract

The annual NBA draft is a constrained fair-allocation problem: 30 indivisible draft picks are assigned to 30 teams that take turns making selections. The top picks are allocated through a lottery for the 14 non-playoff teams. The lottery has a documented manipulation incentive (“tanking”): a team eliminated from playoff contention can improve its expected draft position by losing additional games. Munro and Banchio [2020] prove an impossibility result, showing that no mechanism based on end-of-season records alone can simultaneously reward poor performance and eliminate this incentive. Carry-Over Lottery Allocation (Cola) sidesteps the impossibility by replacing single-season records with a multi-year playoff history as the basis for draft priority [Highley et al., 2026]. The incentive-compatibility argument for Cola relies on five assumptions, four of which hold by construction; the fifth (that pool-manipulation gains are negligible) is supported by simulation in the original paper but lacks a closed-form bound. We make four contributions. First, we derive an analytic upper bound of approximately 4% on the worst-case probability gain from pool manipulation under reasonable distributions of teams’ draft priorities. Second, we present the Cola family as five variants on a shared scaffold. Third, we prove an analytic per-series bound on tanking exposure in the Capped Cola variant. Fourth, we analyze Cola according to the axiomatic draft-rule literature [Coreno and Balbuzanov, 2022].

1 Introduction

In the 2025-26 NBA season, multiple teams were publicly accused of tanking in anticipation of a strong draft class. The accusations did not name a small minority: by some counts, roughly a third of the league’s 30 teams were openly losing games at season’s end to improve their lottery odds. In response, the league introduced a number of potential draft reform ideas, including the

“3-2-1 lottery.” [Highley, 2026b]. The proposal includes a three-year re-evaluation provision, which means the design discussion is expected to reopen in 2029.

The annual NBA draft is a constrained fair-allocation problem. Thirty indivisible items (draft picks) are assigned to 30 agents (teams) in priority order. Picks 15–30 are assigned by reverse regular-season record. Picks 1–14 are subject to a lottery confined to the 14 non-playoff teams, with weighted odds tilted toward teams with worse records. (In the current system, the bottom three teams are weighted equally and the rest declining). Throughout the years, the NBA has adjusted odds while facing the same question: how to assign indivisible items to agents while (a) preferentially rewarding the weakest agents, (b) eliminating the incentive to under-perform, and (c) producing a mechanism transparent enough that fans, owners, and players can reason about it.

Munro and Banchio [2020] give a formal treatment of the central tension and prove an impossibility. They define a No-Tanking Condition (NTC): no team can improve its expected outcome by losing. Any mechanism mapping end-of-season records to draft probabilities, while giving worse-record teams strictly better odds, must violate NTC. Only a uniform lottery (equal odds for all non-playoff teams) satisfies NTC, abandoning differential reward. The 2026 “3-2-1 lottery” proposal takes precisely this route, with adjustments for the play-in tournament and reduced odds for the three teams at the bottom of the standings [Highley, 2026b].

Carry-Over Lottery Allocation (Cola) sidesteps the impossibility by replacing single-season records with multi-year playoff history as the basis for draft priority [Highley et al., 2026]. Each non-playoff team accumulates draft priority over consecutive seasons; success in the playoffs or the draft itself diminishes a team’s draft priority. The single-season manipulation channel that drives the Munro-Banchio counterexample becomes ineffective: no single-season choice materially shifts the multi-year ticket pool. Cola is a family of mechanisms rather than a single design; the original paper develops one concrete instance (Classic Cola), and subsequent work introduces variants that occupy different points in the design space.

Contributions:

86	• (Section 4) An analytic manipulation bound: Cola’s	when IC would first be violated requires precise team-	142
87	incentive-compatibility argument rests on an as-	level valuations of playoff-versus-draft trade-offs that the	143
88	sumption that is supported empirically, but until	league does not have access to. Third, the truncation	144
89	now had no analytic bound. We derive a closed-form	timing assumes teams have an incentive to win at the	145
90	upper bound of approximately 4% on the worst-case	start of the season, but some teams have an incentive to	146
91	probability gain from pool manipulation, under rea-	tank from the first game.	147
92	sonable assumptions about the underlying param-		
93	eters. The bound contains the simulation-observed	2.2 Preference Manipulation and the Axiomatic	148
94	extremes (a 3% maximum reported by Highley et	Question	149
95	al. [2026] across 1,000 synthetic seasons).	Coreno and Balbuzanov [2022] address a complementary	150
96	• (Section 3–4) A family-level perspective: Highley	manipulation channel: teams misreporting which players	151
97	et al. [2026] describe numerous ways that variants	they prefer. They prove that no allocation rule simulta-	152
98	could be introduced to Cola. Here, we present a	neously satisfies:	153
99	scaffold for five concrete variants (Section 3).	• Respect for Priority (RP): higher-priority teams re-	154
100	• (Section 3–4) Bound on cost of winning One variant	ceive weakly better allocations [Svensson, 1994];	155
101	is Capped Cola [Highley, 2026b], which introduces a	• Envy-Free up to One Object (EF1): no team envies	156
102	stockpile cap to bound the per-series cost of advanc-	another’s allocation after removing one item [Lipton	157
103	ing in the playoffs. We prove an analytic per-series	et al., 2004; Budish, 2011];	158
104	bound (Lemma 3): under cap Max, no playoff series	• Strategy-Proofness (SP): no team benefits from mis-	159
105	win ever costs a team more than $0.3 \cdot \text{Max}$ tickets	reporting preferences.	160
106	(45 at $\text{Max} = 150$).	RP extends Svensson’s “weak fairness” from queue al-	161
107	• (Section 2.2) A bridge to the axiomatic litera-	location to draft priority [Svensson, 1994]. EF1 origi-	162
108	ture: Coreno and Balbuzanov [2022] prove that no	ates with Lipton et al. [2004] for indivisible-goods al-	163
109	draft rule simultaneously satisfies Respect for Pri-	location in computer science settings, and is adapted to	164
110	ority (RP), Envy-Free up to One Object (EF1),	the assignment problem by Budish [2011] in economics.	165
111	and Strategy-Proofness (SP). Their result is single-	The Coreno-Balbuzanov result is single-period: it	166
112	period; Cola’s multi-year structure raises an open	characterizes one-shot allocations of indivisible items.	167
113	question about whether the axioms lift. We sur-	Cola is multi-period: a team’s priority in year t is a	168
114	face a preliminary mapping and identify what would	function of its outcomes in years $t - 1, t - 2, \dots$. The lift	169
115	need to be shown to embed Cola in the axiomatic	splits cleanly: the deterministic variant (Simple Cola)	170
116	literature.	inherits the axioms exactly; the lottery-based variants	171
117	A 26-season historical backtest of all five variants (Sec-	(Simple Lottery, Classic, Countdown, Capped) inherit	172
118	tion 5) calibrates the bounds against real NBA data,	them in expectation. The precise named probabilistic	173
119	with the Philadelphia 76ers’ 2013–2017 “Process” as a	counterparts for the lottery case are open.	174
120	case study in how multi-year diminishment dismantles	RP under Cola: Simple Cola assigns picks determinis-	175
121	repeat-tanking strategies that the current weighted lot-	tically by drought (with a regular-season tiebreak) and	176
122	tery rewards.	satisfies RP exactly. The lottery-based variants satisfy	177
123	2 Background and Related Work	RP in expectation: $\mathbb{E}[\text{pick rank} \mid L_i]$ is monotone in L_i ,	178
124	2.1 Effort Manipulation: The Munro-Banchio	though the lottery’s randomization can place a higher-	179
125	Impossibility	priority team behind a lower-priority one in any single	180
126	Munro and Banchio [2020] establish that any mecha-	draw. The natural target is a stochastic-dominance re-	181
127	nism mapping end-of-season records to draft probabil-	finement of Svensson’s priority property; pinning down	182
128	ities, while giving the worst team strictly better odds,	the precise named axiom is open.	183
129	must violate NTC. The proof constructs a minimal coun-	EF1 under Cola: The same dichotomy holds. Sim-	184
130	terexample: two eliminated teams with identical records	ple Cola’s deterministic allocation satisfies EF1 directly:	185
131	and a final game between them. The loser finishes with	with picks as the indivisible items, no team envies an-	186
132	a worse record and therefore receives better draft odds;	other’s allocation after removing one item from the en-	187
133	both teams prefer to lose. Only the uniform lottery sat-	vied bundle. The lottery-based variants satisfy EF1 in	188
134	isfies NTC, but uniform odds abandon the goal of differ-	expectation. A stochastic-dominance form of EF1, in the	189
135	entially rewarding the weakest teams.	spirit of the random-assignment literature, is the natu-	190
136	The authors propose T-IC (Truncated Incentive Com-	ral probabilistic counterpart; the right form for history-	191
137	patibility) as an alternative. T-IC tracks each team’s	based priority rules is open.	192
138	conditional probability of finishing last and freezes allo-	SP under Cola: Coreno-Balbuzanov’s strategy-	193
139	cations at the moment IC would first be violated. The	proofness concerns preference manipulation in draft	194
140	mechanism is optimal in theory but has three practical	selection (whether a team is incentivized to pass on	195
141	issues. First, it is opaque to fans. Second, calculating	its truly preferred player in round one to acquire that	196

player in round two) and is not material to Cola’s design, which targets priority order rather than draft-day pick choice. The strategy-proofness that matters here is resistance to effort manipulation (tanking) and pool manipulation: Section 4 bounds the pool-manipulation channel under Classic Cola, and Lemma 3 bounds the per-series channel under Capped Cola.

3 The COLA Mechanism Family

3.1 Common Framework

Every Cola variant maintains a per-team integer-valued state across seasons and uses that state to determine draft priority in the current season. We refer to this state as the lottery index and write $L_i^{(t)}$ for team i at the end of season t . Depending on the variant, draft priority may be carried over through drought length or a stockpile of lottery tickets. These may be other manifestations of the lottery index.

Definition 1 (Lottery Index). For each team i and season t , the lottery index $L_i^{(t)} \in \mathbb{Z}_{\geq 0}$ is determined by an initial value $L_i^{(0)}$, a per-season increment $\Delta_i^{(t)}$ (variant-specific), and a per-season diminishment $\rho_i^{(t)} \in [0, 1]$ that scales the index down based on playoff and draft outcomes:

$$L_i^{(t+1)} = \left(L_i^{(t)} + \Delta_i^{(t)} \right) \cdot \left(1 - \rho_i^{(t)} \right).$$

The pool in season t is $P^{(t)} = \sum_{i \in E^{(t)}} L_i^{(t)}$, where $E^{(t)}$ is the set of teams eligible for the lottery.

The three components (the eligibility set $E^{(t)}$, the increment $\Delta_i^{(t)}$, and the diminishment $\rho_i^{(t)}$) are the primary levers that distinguish the variants.

3.2 Five Variants

The variants below trade off explainability, robustness, and per-series tanking exposure differently. Table 1 gives a summary of the comparison.

Simple Cola eliminates the lottery: eligibility is the 22 teams with zero playoff series wins (14 non-playoff teams + 8 first-round losers). Teams are ranked by drought (consecutive seasons with neither a playoff series win nor top-3 lottery pick), with ties broken by most regular-season wins. Draft picks are assigned deterministically.

Simple Lottery Cola shares Simple’s drought state and 22-team pool. Pre-2019 NBA lottery odds (25.0% down to 0.5% across the worst 14 teams) are applied to the top 14 by drought; the remaining 8 receive picks 15–22 deterministically.

Classic Cola is the variant developed in detail by Highley et al. [2026], and it is the variant our manipulation bound directly analyzes. The eligibility set is the 14 non-playoff teams. Each non-playoff team receives an increment of $\alpha = 1,000$ tickets per season; playoff teams receive zero increment. Diminishment is graduated: in the playoff channel, champion $\rho = 1.0$, finals 0.75, conference finals 0.5, second round 0.25, first round 0; in

the draft channel, pick #1 $\rho = 1.0$, #2 0.75, #3 0.5, #4 0.25, #5 and below 0. The lottery for picks 1–4 is weighted by lottery index ($p_i = L_i/P$); picks 5–14 are assigned by reverse standings.

Countdown Cola ranks teams by McCarty number $M_i = d_i \cdot w_i$ (drought $\times$ regular-season wins). The same 22-team pool as Simple Cola is dealt with in 5-team elimination pools; draft picks are assigned bottom-up via ticket allocation (2, 3, 4, 5, 6) within each pool. Teams with lower priority have more tickets, and the selected team picks last. A new team is added to the pool, and the process is repeated for the next pick. The mechanism guarantees that no team falls more than four spots below its priority rank.

Capped Cola [Highley, 2026b] introduces a stockpile cap at Max (default 150). The increment is $\Delta_i^{(t)} = w_i^{(t)}$ for any team finishing the regular season outside the top 6 of either conference; eligibility excludes (i) teams that won a top-5 lottery pick last season and (ii) teams that won a playoff series this season or last. Diminishment uses a six-step playoff schedule (champion 1.0, runner-up 0.8, conference finals 0.6, second round 0.4, first round 0.2, play-in winner losing round 1 0.1, play-in loser 0) and a five-step draft schedule (#1 1.0, #2 0.8, #3 0.6, #4 0.4, #5 0.2). The cap smooths the diminishment cliffs of drought length-based versions of Cola and bounds the per-series tanking incentive (Lemma 3).

3.3 Incentive Compatibility

Highley et al. [2026] establish IC for Classic Cola under five assumptions:

1. **Playoff primacy:** the value to a team of advancing in the playoffs exceeds the value of any plausible draft-pick gain.
2. **Playoff advancement:** within the playoffs, advancing further is preferred to losing earlier, controlling for opponent strength.
3. **Picks 5–14 negligible at the margin:** no team profits from manipulating its position among those slots through within-season effort.
4. **Pool manipulation negligible:** no team can meaningfully shift its expected draft position by manipulating which other teams are in or out of the lottery pool.
5. **No traded-pick protections:** pick trades do not interact with diminishment in a way that creates manipulation incentives.

Assumptions (1)–(3) and (5) hold by construction or follow from straightforward arguments. Assumption (4) is supported by Highley et al. [2026] via a 1,000-season simulation, observing a maximum probability gain of approximately 3%. Section 4 converts this empirical observation into a closed-form analytic bound.

Variant	Pool	Increment	Diminishment	Lottery
Simple	22: sw = 0	drought +1	reset on snapped drought	deterministic by drought
Simple Lottery	22: sw = 0	drought +1	reset on snapped drought	pre-2019 NBA odds (top 14), then by drought
Classic	14: non-playoff	+1000 if non-playoffs	graduated, 4 levels	top-4 raffle by index, then by standings
Countdown	22: sw = 0	drought +1	reset on snapped drought	rank by drought $\times$ wins; nested elimination pools
Capped	exclusion-rule	$+w_i^{(t)}$ if play-in or below	graduated, 6/5 levels	top-5 raffle by index, then by index

Table 1: The Cola family. “sw = 0” is shorthand for “zero playoff series wins this season.” Capped Cola’s exclusion-rule pool removes teams that won a top-5 lottery pick last season or a playoff series this/last season. Each variant occupies a distinct point of the design space; the manipulation bound (Theorem 2) is established for Classic and the per-series bound (Lemma 3) for Capped.

4 Pool Manipulation Bound

4.1 Setup

Let $P = \sum_{i=1}^n L_i$ denote the total lottery pool, where L_i is team i ’s lottery index and n is the number of lottery-eligible teams. A team’s probability of winning the first pick is $p_i = L_i/P$.

The pool manipulation channel operates as follows. Team i deliberately loses to a high-index opponent h near the playoff boundary, causing h to make the playoffs. Team h (with index L_h) exits the lottery pool and is replaced by a lower-index team ℓ (with index $L_\ell < L_h$). The pool shrinks by $\Delta = L_h - L_\ell > 0$.

Proposition 1 (Manipulation Gain). The probability gain for team i from a pool swap of size $\Delta = L_h - L_\ell$ is $G_i = L_i \cdot \Delta/[P \cdot (P - \Delta)]$.

Proof. Before manipulation, $p_i = L_i/P$. After the swap, the new pool is $P' = P - \Delta$ and $p'_i = L_i/P'$. The gain is $G_i = p'_i - p_i = L_i(1/(P - \Delta) - 1/P) = L_i \cdot \Delta/[P(P - \Delta)]$. $\square$

When $\Delta \ll P$ (typically $\Delta/P \approx 4\%$ in our backtest), the gain factors into two ratios: the manipulator’s pool share (p_i) and the fractional pool change (Δ/P). Both have structural ceilings, which we exploit below.

Indexing convention. The bound below uses the expression $L_i = T_i \cdot \alpha$, where T_i is team i ’s consecutive non-playoff seasons. For Classic Cola, this is exact for any team starting from a stockpile reset (a championship or top-5 lottery pick) and is the steady-state attractor under the diminishment schedule. A fully general bound that handles ticket inheritance from earlier playoff series under the full diminishment dynamics is in the broader paper.

4.2 Conditioned Upper Bound

Definition 2 (Structural Anti-Correlation). Let $T_{\max}$ denote the maximum consecutive non-playoff years for any team, and T_{boundary} the maximum consecutive non-playoff years for a team near the playoff boundary (a team that could plausibly be pushed into the playoffs by a single loss). The index distribution exhibits structural anti-correlation when $T_{\text{boundary}} \ll T_{\max}$: teams with the highest indices are far from the playoff boundary and thus cannot be the target of pool manipulation.

Theorem 2 (Conditioned Manipulation Bound). Under structural anti-correlation with parameters $T_{\max}$, T_{boundary} , and average index multiplier $\bar{k} = \bar{L}/\alpha$, where α is the annual index increment, the maximum manipulation gain satisfies


$$G_{\max} \leq \frac{T_{\max} \cdot (T_{\text{boundary}} - 1)}{(T_{\max} + n - 1) \cdot n \cdot \bar{k}}.$$

Proof sketch. The gain is maximized when L_i is maximized and Δ/P is maximized. Bounding $p_{\max}$: the highest-index team has $L_{\max} = T_{\max} \cdot \alpha$. The minimum pool size occurs when one team has $T_{\max}$ years of accumulation and the remaining $n - 1$ teams each have the minimum α : $P_{\min} = (T_{\max} + n - 1) \cdot \alpha$. Thus $p_{\max} \leq T_{\max}/(T_{\max} + n - 1)$. Bounding Δ/P : the team being pushed out has at most $T_{\text{boundary}} \cdot \alpha$ tickets, and the team replacing it has at least α , so $\Delta_{\max} = (T_{\text{boundary}} - 1) \cdot \alpha$. In steady state, $P \approx n \cdot \bar{k} \cdot \alpha$. Combining via the first-order approximation yields the stated bound. $\square$

Empirical calibration: We instantiate the bound with $T_{\max} = 10$ (the typical-case ceiling on consecutive non-playoff years across teams in our 26-season backtest), $T_{\text{boundary}} = 5$ (a conservative ceiling on the consecutive non-playoff years a near-boundary team can accumulate before improving enough to clear the boundary), $n = 14$ (the structural count of NBA non-playoff teams), and $\bar{k} = 3.1$ (an average team has roughly three years of carry-over before reset, matching the original paper’s observation). These yield

$$G_{\max} \leq \frac{10 \cdot 4}{23 \cdot 14 \cdot 3.1} \approx 4.0\%,$$

which contains the simulation-observed maximum of 3.0% across 1,000 seasons [Highley et al., 2026]. An independent 200-season Bradley-Terry simulation we ran confirms the bound, with observed maximum gain of 1.70%. The 26-season historical dataset contains one outlier: the Sacramento Kings’ 16-year drought (2006–07 through 2021–22). Setting $T_{\max} = 16$ to include the outlier loosens the bound to approximately 5.1%, still well above the empirical extremes; we report the typical-case bound ($T_{\max} = 10$) as the headline because the next-longest droughts are Minnesota at 13 years and Phoenix and Charlotte at 10 each.



(a) Bound vs. (b) Analytic bound vs. simulation.
 $T_{\max}, T_{\text{boundary}}$.

Figure 1: Manipulation bound. (a) Sensitivity to the structural-anti-correlation parameters; the bound is most sensitive to T_{boundary} , and stays under 5% for $T_{\text{boundary}} \leq 5$. (b) Analytic bound (red) vs. the empirical distribution of manipulation gains across 200 Bradley-Terry-simulated seasons; the bound contains all simulated outcomes with margin.

4.3 Sensitivity, Coalitions, and the Pick-Value Gap

The bound’s parameter sensitivity (Figure 1a) shows that T_{boundary} is the dominant lever: even under pessimistic assumptions ($T_{\text{boundary}} = 7$), the bound stays under 7%, and under realistic conditions ($T_{\text{boundary}} \leq 5$) it stays under 5%. This matters operationally because T_{boundary} is the empirical quantity most amenable to direct measurement.

For completeness, coalitional manipulation (two or more teams coordinating to shrink the pool) does not amplify individual gains. Joint gain for a coalition $\{A, B\}$ is $(L_A + L_B) \cdot \Delta / [P(P - \Delta)]$. Even at 30% joint pool share, the gain remains $\leq 1.3\%$, and per-member gain cannot exceed the highest-index individual’s gain. Coordination costs (detection risk, trust problems, sequential uncertainty) further dampen the channel.

The bound addresses the probability term of the manipulation incentive. The value term (the pick-value gap $d_k - d_{k+1}$ between adjacent draft positions) has no formal ceiling in the original specification, though it is bounded for any realistic draft class. Closing assumption (4) entirely would also require a separate argument capping $d_k - d_{k+1}$ via the distribution of draft-eligible prospect quality. We leave this to future work.

4.4 Per-Series Bound for Capped Cola

Capped Cola introduces a fourth manipulation channel: tanking through the playoffs themselves (a team within the cap declining to advance, in order to retain its stockpile). The cap converts the playoff-diminishment schedule from a fraction-of-current-stockpile cost (unbounded above when stockpile grows large) to a fraction-of-Max cost (bounded by construction).

Lemma 3 (Per-Series Tanking Cost). Let Max be the Capped Cola stockpile cap. For any team entering the playoffs with $L_i \leq \text{Max}$, the maximum reduction in L_i attributable to winning (rather than losing) any single

playoff series is bounded by:

$$\Delta L_{\text{series}}^{\max} = \begin{cases} 0.3 \cdot \text{Max} & \text{play-in advancers in round 1,} \\ 0.2 \cdot \text{Max} & \text{any other series advancement.} \end{cases}$$

At Max = 150, the bound is 45 tickets for the play-in case and 30 tickets for any subsequent series.

Proof. Reading off Capped Cola’s six-step diminishment schedule, the per-series increments in the diminishment fraction $\Delta\rho$ are: play-in advancer round 1 ($\rho_{\text{lose}} = 0.1$, $\rho_{\text{advance}} = 0.4$, $\Delta\rho = 0.3$); standard round 1 ($0.2 \rightarrow 0.4$, $\Delta\rho = 0.2$); round 2 to conference finals ($0.4 \rightarrow 0.6$); conference finals to NBA finals ($0.6 \rightarrow 0.8$); finals to championship ($0.8 \rightarrow 1.0$). All non-play-in transitions yield $\Delta\rho = 0.2$. The cost in tickets is $\Delta\rho \cdot L_i^{\text{pre}} \leq \Delta\rho \cdot \text{Max}$, with equality at the cap. $\square$

The bound makes the playoff-tanking incentive transparent at the league level: at Max = 150, no single series win ever costs a team more than 45 lottery tickets, which is approximately a year and a half of accumulation for a team near the wins-increment ceiling. The cumulative regret for a finals-losing team at the cap is $0.8 \cdot \text{Max} = 120$ tickets, decomposed into per-series increments per Lemma 3.

4.5 Three Closures of Assumption (4)

The pool-manipulation gap (assumption (4) in Section 3) admits at least three closure strategies, each represented in the literature.

- Empirical closure: Highley et al. [2026] simulate 1,000 synthetic seasons under Bradley-Terry and observe a maximum gain of $\sim 3\%$. The closure is statistical: the assumption holds in the typical case but admits no parameter-explicit ceiling.
- Behavioral closure: the companion formal-proof video to [Highley et al., 2026] introduces an additional behavioral assumption $G > \text{expected pool-manipulation gain}$ (where G is the value of a single regular-season win), eliminating the channel via preference structure rather than mechanism design.
- Analytic closure (this paper): Theorem 2 provides a closed-form upper bound conditioned on structural properties of the index distribution. The bound is parameter-explicit and admits sensitivity analysis (Figure 1a).

The three closures compose: any practical Cola deployment can rely on whichever closure best matches the available evidence. Our contribution is the analytic option, which previously did not exist.

5 Empirical Validation

5.1 Backtesting Methodology

We compute Cola state for every team across 26 NBA seasons (1999–2025) using verified Basketball Reference data: 775 team-season records, 776 first-round draft picks, and 55 lottery-day-attributed pre-draft pick

Season	W-L	Pick	Cola index	Rank
2013-14	19-63	#3	2,750	8/14
2014-15	18-64	#3	2,375	9/14
2015-16	10-72	#1	2,188	10/14
2016-17	28-54	#3	1,000	14/14

Table 2: Philadelphia 76ers under Classic Cola during the “Process.” Each top-3 pick triggers graduated diminishment (#1: full reset; #3: 50% reduction), leaving Philadelphia at the bottom of the lottery pool by the fourth tanking season despite a 10-72 record. The Sacramento Kings hold the top Classic Cola index throughout this period, climbing from 7,250 in 2013-14 to 10,250 in 2016-17 via organic non-competitiveness.

trades. The engine separates pre-draft state (Phase A, used for draft-probability calculations) from post-draft diminishment (Phase B, carried forward to the next season). Cold-start effects stabilize by approximately season 5; we report aggregates over 21 stable seasons (2005-2025).

Two methodological points warrant attention. First, an audit of all 364 lottery picks (14 picks $\times$ 26 seasons) identified 8 instances where the draft-day holder differed from the lottery-day holder via post-lottery trades. Cola requires lottery-day attribution: the team that held the pick when the lottery ran absorbs the diminishment, regardless of who selects. Second, the backtest is static: it applies Cola rules to actual historical outcomes. Under a real Cola regime, draft outcomes would differ in ways that change future drought lengths and indices. We caveat the comparative cross-variant claims accordingly.

5.2 Case Study: The Philadelphia “Process”

The 2013-14 through 2016-17 Philadelphia 76ers represent the canonical strategic-tanking sequence in modern NBA history: four consecutive top-3 lottery outcomes (picks #3, #3, #1, #3) acquired through deliberate multi-season non-competitiveness. Under the current weighted lottery, each tanked season yielded a top-tier pick. Under Classic Cola, graduated diminishment makes the same sequence self-defeating. Table 2 shows Philadelphia’s Classic Cola trajectory.

Two properties emerge. First, diminishing returns to repeat tanking: the Classic Cola diminishment schedule ($\{1.0, 0.75, 0.5, 0.25\}$ for picks #1-4) dominates the $\alpha = 1,000$ annual accumulation, so a team that receives a top-3 pick in year t cannot re-enter the top of the lottery in year $t + 1$ through tanking alone. Second, pick attribution matters: the comparison is sensitive to who holds the pick on lottery day. The 2017 Boston-Philadelphia Fultz/Tatum swap was a post-lottery trade; under our convention, Boston’s index absorbs the #1 diminishment, not Philadelphia’s.

5.3 Capped Cola Cap Calibration

Capped Cola is parameterized by a single tunable constant Max. The cap controls the trade-off between (i) ordering resolution at the top of the lottery (a higher cap

permits more separation between the highest-stockpile teams) and (ii) per-series tanking exposure (Lemma 3: a higher cap raises the cost of advancing past round 1). Table 3 reports 21-stable-season aggregates from a sweep over $\text{Max} \in \{75, 100, 125, 150, 175, 200\}$.

Three patterns emerge. Sub-100 caps over-compress the lottery: at $\text{Max} = 75$, more than six teams sit at the cap on average and the rank-1 vs. rank-5 separation collapses to under one ticket, destroying the priority signal the family inherits from Classic. Caps above 175 erase the cap’s purpose: fewer than one team sits at the cap on average, and the per-series exposure approaches the unbounded behavior of Classic. The published $\text{Max} = 150$ choice produces a mean of 1.62 teams at the cap across the 21 stable seasons (range 0-3), matching the stated headline of “no more than three teams at the maximum at a time, and on average fewer than two” [Highley, 2026b].

5.4 A Note on Trade Handling

A practical question for any deployment is how diminishment interacts with traded picks. Highley [2026a] identifies four options: (1) traded picks affect neither team’s index; (2) the original holder’s index absorbs the diminishment; (3) the receiving team’s index absorbs the diminishment; (4) the diminishment is split. Each option has different incentive properties. Option (2) preserves the link between record and diminishment but over-penalizes teams that traded picks years before knowing they would be top-tier (e.g., the 2011 Clippers, who traded a future first that eventually became Cleveland’s Kyrie Irving pick). Option (3) rewards the wrong actor when the recipient did not earn the pick through their own record. Highley [2026a] recommends option (2) paired with a strengthened Stepien Rule (a cap on the number of future firsts a team can trade), which limits the frequency at which option (2)’s over-penalization surfaces. We do not extend the analysis here; we flag it as a practical implementation question that any Cola deployment will face.

6 Discussion and Open Directions

The bounds in this paper formalize the manipulation-resistance properties of the Cola family. Theorem 2 bounds the open assumption (4) in Classic Cola’s incentive-compatibility argument at approximately 4% under empirically calibrated parameters, with explicit sensitivity to T_{max} , T_{boundary} , n , and $\bar{k}$. Lemma 3 bounds the per-series exposure that Capped Cola’s stockpile cap was designed to address. Combined, the four manipulation channels available under Capped Cola (single-season tanking, pool manipulation, multi-year strategic tanking, and per-series playoff tanking) compose to a small joint exposure relative to the status-quo lottery, where repeat tanking has memoryless returns.

The 2026 “3-2-1” proposal [Highley, 2026b] approximates a uniform lottery within an adjustment for the play-in tournament. Its three-year re-evaluation provision is expected to reopen the design discussion in

Max	Teams at cap (mean)	Teams at cap (max)	Sep. R1–R5 (mean)	Per-series cost (max)	Cum. regret (max)
75	6.71	9	0.9	22.5	45.0
100	4.33	8	11.7	30.0	53.7
125	2.86	4	27.9	37.5	53.7
150	1.62	3	47.7	45.0	57.6
175	0.95	2	65.4	52.5	67.2
200	0.52	2	76.4	60.0	76.8

Table 3: Capped Cola aggregates across 21 stable seasons (2005–2025). Per-series cost max = $0.3 \cdot \text{Max}$ (Lemma 3); cumulative regret max is the worst empirical observation. The bolded Max = 150 row is the variant’s published default [Highley, 2026b].

2029. Cola’s contribution to that re-evaluation is the multi-year-history dimension that 3-2-1 omits, and a family-level view that gives policymakers a continuum across the variants. The bounds in this paper formalize the manipulation-resistance properties any future re-evaluation would need to address.

Open directions: Three follow-on questions extend this work.

- Tighter manipulation bound: The structural anti-correlation in Definition 2 is a calibrated assumption rather than a formally derived consequence of the mechanism. A stationarity argument that bounds T_{boundary} as a function of league parameters (number of teams, season length, playoff structure) would close assumption (4) entirely.
- Pick-value gap: Combining the probability bound (this paper) with a separate bound on the pick-value gap $d_k - d_{k+1}$, derived from the distribution of draft-eligible prospect quality, would close the manipulation incentive in absolute terms, not just in probability.
- Axiomatic embedding: The lift from Coreno and Balbuzanov [2022]’s axioms (RP, EF1, SP) to multi-period priority mechanisms remains open. Section 2.2 sketches the mapping; a formal verification would embed Cola in the existing axiomatic literature on assignment mechanisms and clarify which one-period impossibilities transfer to multi-period settings. We see this as the most natural follow-on for the workshop’s audience.

A reproducibility artifact accompanies this paper: an interactive backtester implementing all five Cola variants plus the 3-2-1 proposal across 26 NBA seasons, with locked-seed Monte Carlo for Countdown Cola probabilities [Rajan, 2026].

References

- Eric Budish. The combinatorial assignment problem: Approximate competitive equilibrium from equal incomes. *Journal of Political Economy*, 119(6):1061–1103, 2011.
- Julian Coreno and Ivan Balbuzanov. Axiomatic characterizations of draft rules. *arXiv preprint arXiv:2204.08300*, 2022.

Timothy Highley, Tannah Duncan, and Ilia Volkov. Carry-over lottery allocation: Practical incentive-compatible drafts. *arXiv preprint arXiv:2602.02487*, 2026.

Timothy Highley. NBA tanking is solvable: Bring your doubts, 2026. <https://highleytj.substack.com/p/nba-tanking-is-solvable-bring-your>.

Timothy Highley. Overcoming inertia to stop NBA tanking. Substack post introducing Capped Cola, 2026. <https://highleytj.substack.com/p/overcoming-inertia-to-stop-nba-tanking>.

Richard J. Lipton, Evangelos Markakis, Elchanan Mossel, and Amin Saberi. On approximately fair allocations of indivisible goods. In *Proceedings of the 5th ACM Conference on Electronic Commerce (EC)*, pages 125–131, 2004.

Evan Munro and Martino Banchio. A no-tanking draft allocation policy. In *MIT Sloan Sports Analytics Conference*, 2020.

Kaushik Rajan. COLA explorer: Interactive backtester for the carry-over lottery allocation family, 2026. <https://kvr06-ai.github.io/cola-manipulation-bound/>; source: <https://github.com/kvr06-ai/cola-manipulation-bound>.

Lars-Gunnar Svensson. Queue allocation of indivisible goods. *Social Choice and Welfare*, 11(4), 1994.